

# CPSC 340 – Tutorial 3

Slides courtesy of Nam Hee Kim

University of British Columbia

# Agenda

---

- Ensemble Methods
  - Bootstrap Aggregating (*Bagging*)
  - Boosting
  - Averaging
  - Random Forest
- Unsupervised learning: Clustering
  - KMeans
  - DBSCAN

# Ensemble methods

---

- Idea: use classifiers as building blocks for more complex model
- Goal: reduce bias/variance by combining several base models

*Reminder:* bootstrapping = random sampling with replacement

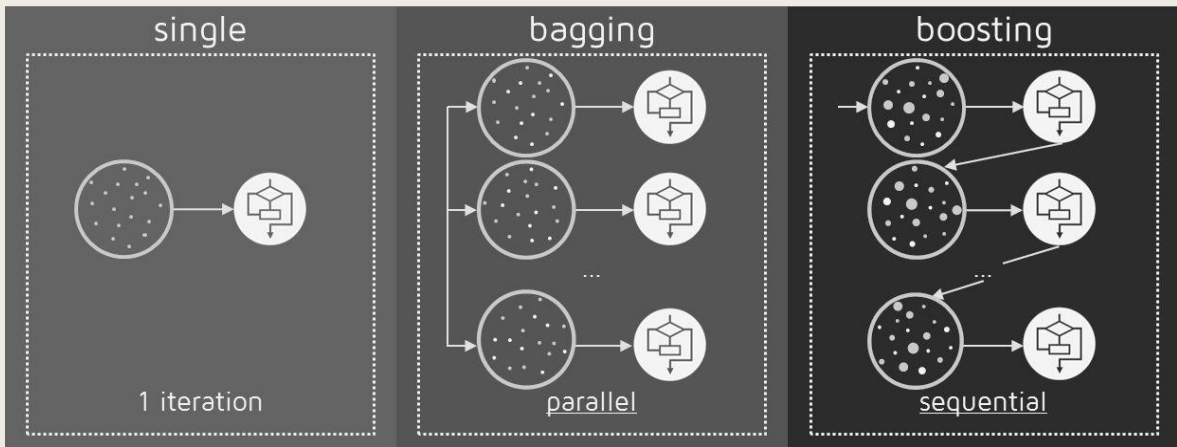
# Ensemble methods: bagging, boosting and averaging

---

- **Bootstrap Aggregating** (bagging): base models are trained **independently** on bootstrap samples, and then combined through majority voting or averaging
- **Boosting**: base models are trained **sequentially**, and then combined through majority voting or averaging. In boosting, we perform “weighted” resampling, so observations that were misclassified by the previously trained base model have a greater chance of being sampled
- **Averaging**: base models are trained **independently**, the final predictions are determined through majority voting or averaging. **Stacking** is a variation of this method in which base models’ predictions serve as features for a *meta-classifier*

# Ensemble methods: bagging vs. boosting

---




# Ensemble methods: bagging vs. boosting

## Quick Questions

---

1. What do we aim to reduce with bagging, bias or variance?
2. How about in boosting?




# Ensemble methods: bagging vs. boosting

## Quick Questions

---

1. What do we aim to reduce with bagging, bias or variance?
  - a. reduce variance, which helps controlling overfitting
2. How about in boosting?
  - a. both, boosting tries to improve the training error for classifiers with high  $E_{\text{train}}$



# Random Forest

---

Idea: bagging using random trees as base classifiers

Random trees: For each tree, randomly sample a small number of features to be considered as candidates split variables.

Given  $n$  (# of training examples),  $d$  (# of features),  $m$  (depth), and  $t$  (# of test examples):

- a) What is the cost of fitting a random tree that uses  $\sqrt{d}$  random features?
- b) What is the cost of predicting using a random tree?



# Random Forest

---

Training:

- Our normal cost for fitting a decision tree is  $O(mnd \log(n))$ , and the factor of  $d$  comes from searching through all features for each split. If we only search  $\sqrt{d}$  features this will be reduced to  $O(mn\sqrt{d} \log(n))$ .

Predicting:

- For each training example  $t$ , we only do an  $O(1)$  operation at each level of the decision tree to decide whether we satisfy the rule. This gives a total cost of  $O(tm)$ .

# Unsupervised Learning

---

- Only have  $x_i$  values, but no explicit target labels.
- **Clustering:** trying to discover the underlying structures (clusters) in the data, such as grouping customers by purchasing behavior.
- **Association:** trying to identify rules that describe large portions of the data, such as people that buy X also tend to buy Y.

# K-Means

---

Idea: group points based on how similar they are to the “mean” of the cluster  
Hyperparameter:  $k$  number of clusters

## Fit:

1. Initialize  $k$  clusters by creating  $k$  centroids (usually randomly)
2. Assign each data point to closest mean
3. Update the means bases on the assignment
4. Repeat until convergence (for example, when points don't change clusters)

## Predict:

1. Assign test example to the nearest mean

## K-Means: Quick questions

---

- What typically happens with  $E_{train}$  as  $k$  increases?
- What is the space complexity for K-Means?

# K-Means: Quick questions

---

What typically happens with  $E_{train}$  as  $k$  increases?

- $E_{train}$  should decrease

What is the space complexity for K-Means?

- $O(kd)$ . You need to store  $k$  cluster centres, each is a vector of length  $d$

# K-Means: Quick questions

---

What typically happens with  $E_{train}$  as  $k$  increases?

- $E_{train}$  should decrease

What is the space complexity for K-Means?

- $O(kd)$ . You need to store  $k$  cluster centres, each is a vector of length  $d$

# K-Means: Practice question

---

Suppose we fit an unsupervised k-means model to a dataset and obtained the following means:

$$W = \begin{bmatrix} 8.0 & 9.5 \\ 1.0 & 7.0 \\ 13.5 & 3.0 \end{bmatrix}$$

Which cluster would a test example  $x = [9 \ 4]$  get assigned to?

# K-Means: Practice question

---

Suppose we fit an unsupervised k-means model to a dataset and obtained the following means:

$$W = \begin{bmatrix} 8.0 & 9.5 \\ 1.0 & 7.0 \\ 13.5 & 3.0 \end{bmatrix}$$

Which cluster would a test example  $x = [9 \ 4]$  get assigned to?

$$(8-9)^2 + (9.5-4)^2 = 31.25$$

$$(1-9)^2 + (7-4)^2 = 73$$

$$(13.5-9)^2 + (3-4)^2 = 21.25$$



# Density-based Spatial Clustering (DBSCAN)

---

Idea: clusters are defined by dense regions (non-dense regions don't get clustered), merge all neighboring core points to form clusters

Hyperparameters:  $\epsilon$  - threshold used to determine if another point is a neighbor

minNeighbors - number of neighbors needed to say that a region is dense enough to be a cluster.

"Core" points have at least minNeighbors

## DBSCAN: Quick questions (T/F)

---

1. For data points to be in a cluster, they must be in a distance threshold to a core point
2. It has strong assumptions for the distribution of data points in dataspace
3. It is a non-parametric model

# DBSCAN: Quick questions (T/F)

---

1. For data points to be in a cluster, they must be in a distance threshold to a core point
  - True
2. It has strong assumptions for the distribution of data points in dataspace
  - False, DBSCAN can form a cluster of any arbitrary shape and does not have strong assumptions for the spatial distribution of data points
3. It is a non-parametric model
  - True

# Bonus Questions

---

1. When choosing  $k$  in K-Means, why not just choose the  $k$  that leads to the smallest distance (sum of squared distances within clusters)?
2. You decide to use clustering for outlier detection; that is, to detect instances that are very atypical compared to all the rest. How might you do this with K-Means?
3. You decide to use clustering for outlier detection; that is, to detect instances that are very atypical compared to all the rest. How might you do this with DBSCAN?